

# linear-ecom-customers

March 19, 2024

```
[1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv("ecomm-customers.csv")
df
```

```
[2]:
```

	Email \
0	mstephenson@fernandez.com
1	hduke@hotmail.com
2	pallen@yahoo.com
3	riverarebecca@gmail.com
4	mstephens@davidson-herman.com
..	...
495	lewisjessica@craig-evans.com
496	katrina56@gmail.com
497	dale88@hotmail.com
498	cwilson@hotmail.com
499	hannahwilson@davidson.com

	Address	Avatar \
0	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine
..	...	...
495	4483 Jones Motorway Suite 872\nLake Jamiefurt,...	Tan
496	172 Owen Divide Suite 497\nWest Richard, CA 19320	PaleVioletRed
497	0787 Andrews Ranch Apt. 633\nSouth Chadburgh, ...	Cornsilk
498	680 Jennifer Lodge Apt. 808\nBrendacheater, TX...	Teal
499	49791 Rachel Heights Apt. 898\nEast Drewboroug...	DarkMagenta

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179

4	33.330673	12.795189	37.536653	4.446308
..	...	...	...	...
495	33.237660	13.566160	36.417985	3.746573
496	34.702529	11.695736	37.190268	3.576526
497	32.646777	11.499409	38.332576	4.958264
498	33.322501	12.391423	36.840086	2.336485
499	33.715981	12.418808	35.771016	2.735160

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092
..	...
495	573.847438
496	529.049004
497	551.620145
498	456.469510
499	497.778642

[500 rows x 8 columns]

```
[3]: df.shape
```

```
[3]: (500, 8)
```

```
[4]: df.isnull().sum()
```

```
[4]: Email          0
      Address        0
      Avatar         0
      Avg. Session Length  0
      Time on App      0
      Time on Website  0
      Length of Membership  0
      Yearly Amount Spent  0
      dtype: int64
```

```
[5]: x = df.drop(columns=["Email", "Address", "Avatar", "Yearly Amount_
      ↪Spent"],axis=1)
      x
```

	Avg. Session Length	Time on App	Time on Website	Length of Membership
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543

3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308
..	...	...	...	...
495	33.237660	13.566160	36.417985	3.746573
496	34.702529	11.695736	37.190268	3.576526
497	32.646777	11.499409	38.332576	4.958264
498	33.322501	12.391423	36.840086	2.336485
499	33.715981	12.418808	35.771016	2.735160

[500 rows x 4 columns]

```
[6]: y = df["Yearly Amount Spent"]
      y
```

```
[6]: 0      587.951054
      1      392.204933
      2      487.547505
      3      581.852344
      4      599.406092
      ...
      495     573.847438
      496     529.049004
      497     551.620145
      498     456.469510
      499     497.778642
      Name: Yearly Amount Spent, Length: 500, dtype: float64
```

```
[7]: from sklearn.model_selection import train_test_split
```

```
[8]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42)
```

```
[9]: from sklearn.linear_model import LinearRegression
```

```
[10]: model = LinearRegression().fit(x_train, y_train)
      model
```

```
[10]: LinearRegression()
```

```
[11]: y_pred = model.predict(x_test)
      y_pred
```

```
[11]: array([402.86230051, 542.53325708, 426.62011918, 501.91386363,
      409.6666551 , 569.92155038, 531.50423529, 505.94309188,
      408.10378607, 473.45942928, 441.18668812, 424.52463471,
      424.83341694, 527.12061508, 430.87985533, 423.47062047,
      575.8751518 , 484.6563331 , 457.77896975, 481.58742311,
      501.56110993, 513.12815188, 507.49166899, 646.63377343,
```

```

449.70050586, 496.26290484, 556.18523776, 554.78684161,
399.1582784 , 325.16921284, 532.62732659, 477.73025415,
500.76491535, 305.09971374, 505.46811902, 483.52069444,
519.09464122, 437.75549737, 456.25005245, 470.63517876,
494.11207805, 444.65549239, 508.57079732, 500.88197484,
488.35128728, 535.34025218, 594.58301773, 513.59474408,
279.69877702, 432.71590835, 421.06976164, 480.94327496,
584.59481888, 608.61734059, 564.42312991, 494.47224504,
393.95593318, 456.11321352, 572.92228417, 499.27385693,
512.42973545, 391.56170305, 479.60705887, 481.05023229,
474.71926117, 546.37716047, 430.11675694, 601.91418143,
422.26508516, 493.11622454, 528.10614863, 581.06630842,
620.60774498, 512.47838603, 411.2147464 , 498.07095351,
461.44587681, 445.63453258, 447.63898998, 534.81030495,
598.85091016, 619.46554961, 494.43362232, 672.2442837 ,
532.15516513, 438.41740681, 514.80907179, 546.73893548,
331.73069072, 510.33949236, 536.21660556, 499.50696031,
375.86919792, 573.61952185, 479.18212334, 588.32862943,
485.18137257, 455.93070091, 398.67820721, 451.70869105])

```

```
[12]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
[13]: r2_sc = r2_score(y_test, y_pred)
```

```
[14]: print(f"R2 Score = ", r2_sc)
```

R2 Score = 0.9778130629184127

OPTIMIZATION

```
[15]: from sklearn.model_selection import GridSearchCV
```

```
[16]: model = LinearRegression()
model
```

```
[16]: LinearRegression()
```

```
[17]: param_grid = {
    'fit_intercept': [True, False],
    'copy_X': [True, False],
    'n_jobs': [-1, None],
    'positive': [False, True]
}
```

```
[18]: grid_search = GridSearchCV(model, param_grid, cv=5, n_jobs=-1)
grid_search.fit(x_train, y_train)
```

```
[18]: GridSearchCV(cv=5, estimator=LinearRegression(), n_jobs=-1,
                param_grid={'copy_X': [True, False],
                            'fit_intercept': [True, False], 'n_jobs': [-1, None],
                            'positive': [False, True]})
```

```
[19]: best_params = grid_search.best_params_
print("Best Parameters :", best_params)
```

```
Best Parameters : {'copy_X': True, 'fit_intercept': True, 'n_jobs': -1,
'positive': False}
```

```
[20]: best_model = LinearRegression(**best_params)
best_model.fit(x_train, y_train)
best_model
```

```
[20]: LinearRegression(n_jobs=-1)
```

```
[21]: y_pred = best_model.predict(x_test)
y_pred
```

```
[21]: array([402.86230051, 542.53325708, 426.62011918, 501.91386363,
409.6666551 , 569.92155038, 531.50423529, 505.94309188,
408.10378607, 473.45942928, 441.18668812, 424.52463471,
424.83341694, 527.12061508, 430.87985533, 423.47062047,
575.8751518 , 484.6563331 , 457.77896975, 481.58742311,
501.56110993, 513.12815188, 507.49166899, 646.63377343,
449.70050586, 496.26290484, 556.18523776, 554.78684161,
399.1582784 , 325.16921284, 532.62732659, 477.73025415,
500.76491535, 305.09971374, 505.46811902, 483.52069444,
519.09464122, 437.75549737, 456.25005245, 470.63517876,
494.11207805, 444.65549239, 508.57079732, 500.88197484,
488.35128728, 535.34025218, 594.58301773, 513.59474408,
279.69877702, 432.71590835, 421.06976164, 480.94327496,
584.59481888, 608.61734059, 564.42312991, 494.47224504,
393.95593318, 456.11321352, 572.92228417, 499.27385693,
512.42973545, 391.56170305, 479.60705887, 481.05023229,
474.71926117, 546.37716047, 430.11675694, 601.91418143,
422.26508516, 493.11622454, 528.10614863, 581.06630842,
620.60774498, 512.47838603, 411.2147464 , 498.07095351,
461.44587681, 445.63453258, 447.63898998, 534.81030495,
598.85091016, 619.46554961, 494.43362232, 672.2442837 ,
532.15516513, 438.41740681, 514.80907179, 546.73893548,
331.73069072, 510.33949236, 536.21660556, 499.50696031,
375.86919792, 573.61952185, 479.18212334, 588.32862943,
485.18137257, 455.93070091, 398.67820721, 451.70869105])
```

```
[22]: r2_sc = r2_score(y_test, y_pred)
```

```
[23]: print(f"R2 Score = ", r2_sc)
      print("Best Parameters :", best_params)
```

R2 Score = 0.9778130629184127

Best Parameters : {'copy\_X': True, 'fit\_intercept': True, 'n\_jobs': -1,  
'positive': False}

```
[ ]:
```