

# logistic-heart

March 19, 2024

```
[1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv("heart.csv")
df
```

```
[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..	...	...	..	...	...	...	...	...	...	...	
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..	...	..	...	...
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

```
[ ]:
```

```
[3]: df.shape
```

```
[3]: (303, 14)
```

```
[4]: df.isnull().sum()
```

```
[4]: age          0
     sex          0
     cp          0
     trestbps     0
     chol         0
     fbs          0
     restecg      0
     thalach      0
     exang        0
     oldpeak      0
     slope        0
     ca           0
     thal         0
     target       0
     dtype: int64
```

```
[5]: df['target'].value_counts()
```

```
[5]: target
     1    165
     0    138
     Name: count, dtype: int64
```

```
[6]: x = df.drop("target", axis=1)
     x
```

```
[6]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..	...	...	..	...	...	...	...	...	...	...	
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

  

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2

```

4          2  0    2
..      ...  ..   ...
298        1  0    3
299        1  0    3
300        1  2    3
301        1  1    3
302        1  1    2

```

[303 rows x 13 columns]

```
[7]: y = df["target"]
```

```
[8]: y
```

```

[8]: 0      1
     1      1
     2      1
     3      1
     4      1
     ..
298    0
299    0
300    0
301    0
302    0
Name: target, Length: 303, dtype: int64

```

```
[9]: df.isnull().sum()
```

```

[9]: age      0
     sex      0
     cp       0
     trestbps  0
     chol     0
     fbs      0
     restecg   0
     thalach   0
     exang     0
     oldpeak   0
     slope     0
     ca        0
     thal      0
     target    0
     dtype: int64

```

```
[10]: from sklearn.model_selection import train_test_split
```

```
[11]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳random_state=42, stratify=y)
```

```
[12]: from sklearn.linear_model import LogisticRegression
```

```
[13]: model = LogisticRegression(max_iter=1000).fit(x_train, y_train)
model
```

```
[13]: LogisticRegression(max_iter=1000)
```

```
[14]: y_pred = model.predict(x_test)
y_pred
```

```
[14]: array([0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,
1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,
0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1], dtype=int64)
```

```
[15]: from sklearn.metrics import accuracy_score, mean_absolute_error,
↳mean_squared_error, r2_score
```

```
[16]: accuracy = accuracy_score(y_test, y_pred)
```

```
[17]: print(f"Accuracy = ", accuracy)
```

Accuracy = 0.8032786885245902

OPTIMIZATION

```
[18]: from sklearn.model_selection import GridSearchCV
```

```
[19]: model = LogisticRegression()
model
```

```
[19]: LogisticRegression()
```

```
[20]: param_grid = {
    'penalty': ['l2', None],
    'solver': ['liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'],
    'C': [1.0, 1.5]
}
```

```
[21]: grid_search = GridSearchCV(model, param_grid, cv=5, n_jobs=-1)
grid_search.fit(x_train, y_train)
```

```
[21]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
    param_grid={'C': [1.0, 1.5], 'penalty': ['l2', None],
    'solver': ['liblinear', 'newton-cg', 'newton-cholesky',
    'sag', 'saga']})
```

```
[23]: best_params = grid_search.best_params_  
print("Best Parameters :", best_params)
```

```
Best Parameters : {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}
```

```
[24]: best_model = LogisticRegression(**best_params)  
best_model.fit(x_train, y_train)  
best_model
```

```
[24]: LogisticRegression(solver='newton-cg')
```

```
[25]: y_pred = best_model.predict(x_test)  
y_pred
```

```
[25]: array([0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,  
          1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,  
          0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1], dtype=int64)
```

```
[26]: accuracy = accuracy_score(y_test, y_pred)
```

```
[27]: print("Best Parameters :", best_params)  
print(f"Accuracy = ", accuracy)
```

```
Best Parameters : {'C': 1.0, 'penalty': 'l2', 'solver': 'newton-cg'}  
Accuracy =  0.8032786885245902
```

```
[ ]:
```

```
[ ]:
```